



 [ece5300sp21](#) / [ece5300](#) Private

ECE 5300 Sp21

 4 stars  0 forks

 Star

 Watch ▼

Code

Issues

Pull requests

Actions

Projects

Wiki



Security


Insights



 master ▼

...




[ece5300](#) / [Demos](#) / [demo05a_logistic_regression.ipynb](#)

 **schniter** updated demo5 (split into demo5a and 5b) and lect5 slides 

 1 contributor

RawBlame

436 lines (436 sloc) | 82.8 KB

Demo: Binary Logistic Regression in One Dimension

```
In [1]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
from sklearn import datasets, linear_model, preprocessing
%matplotlib inline
```

The Logistic Model

This demo helps to visualize an important classification method known as *binary logistic regression*. Logistic regression does not model the binary label $y \in \{-1, 1\}$ as a deterministic function of the features $\mathbf{x} = [x_1, \dots, x_d]^T$, nor as a deterministic function of \mathbf{x} plus additive random noise like we did with linear regression. Instead, it models y as random with a *probability* that is determined by \mathbf{x} through

$$Pr\{y = 1|\mathbf{x}\}.$$

In particular, the logistic model says that this probability is completely determined by the linear *score* or *discriminant* z :

$$Pr\{y = 1|\mathbf{x}\} = f(z), \quad z = b + w_1x_1 + \dots + w_dx_d, \quad (1)$$

(2)

(3)

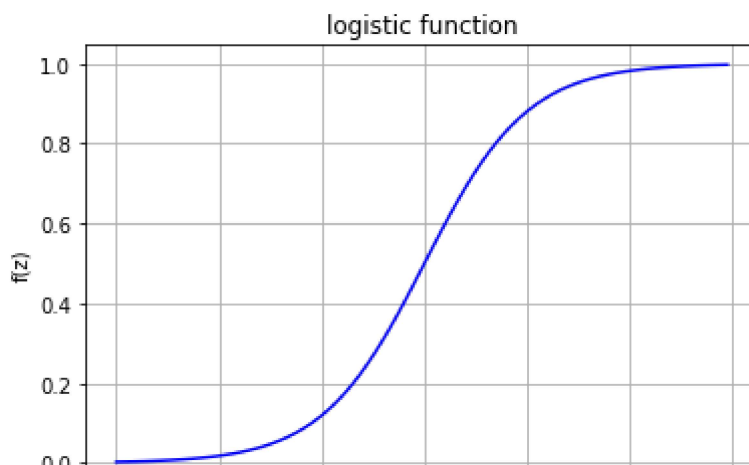
where b is the intercept term (or *bias*), w_1, \dots, w_d are the classifier weights, and $f(\cdot)$ is

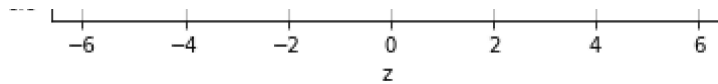
$$f(z) = \frac{1}{1 + e^{-z}},$$

which is known as the *logistic function* or the *sigmoid*. We plot $f(z)$ below.

```
In [2]: def f(z):
return 1/(1+np.exp(-z))

z = np.arange(-6, 6, 0.1)
plt.plot(z, f(z), 'b-')
plt.xlabel('z')
plt.ylabel('f(z)')
plt.title('logistic function')
plt.grid();
```





To illustrate the workings of the logistic function, this demo considers the simple case that $d = 1$ (i.e., one scalar feature). In this case, we can write

$$z = b + wx.$$

We'd like to understand the roles of b and w , as well as how to fit these parameters to data.

Effect of weight and intercept

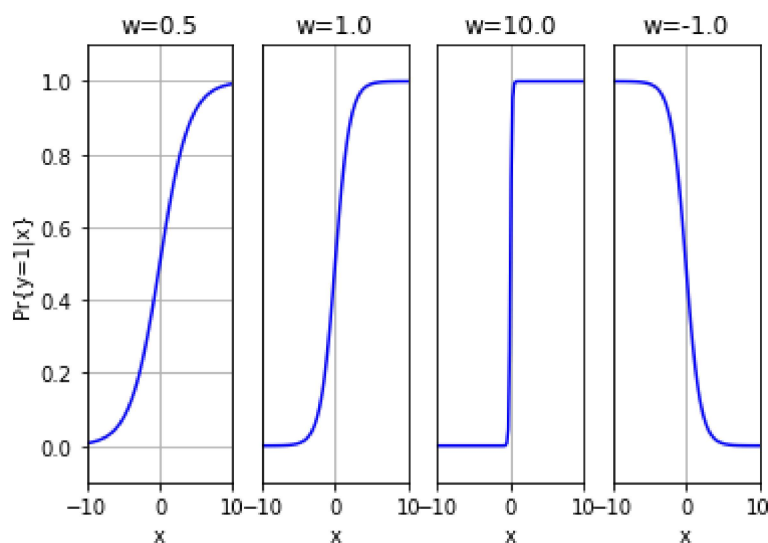
To understand the effect of the weight w , we first set the intercept to zero (i.e., $b = 0$). Below, we plot $\Pr\{y = 1|x\} = f(wx)$ versus x for three different values of w .

```
In [3]: nx = 100
        xmax = 10
        x = np.linspace(-xmax, xmax, 100)

        wls = np.array([0.5, 1, 10, -1])
        nplot = wls.size

        for iplot, w1 in enumerate(wls):
            py = f(w1*x)

            plt.subplot(1, nplot, iplot+1)
            plt.plot(x, py, 'b-')
            plt.axis([-xmax, xmax, -0.1, 1.1])
            plt.grid()
            if (iplot == 0):
                plt.ylabel('Pr{y=1|x}')
            else:
                plt.yticks([])
            plt.xticks([-10, 0, 10])
            plt.xlabel('x')
            plt.title('w={0:.1f}'.format(w1))
```



The figures above show that, when w is positive, $\Pr\{y = 1|x\}$ transitions monotonically from low to high as x increases. Meanwhile, when w is negative, $\Pr\{y = 1|x\}$ transitions monotonically from high to low as x increases. Furthermore, the speed of the transition is proportional to the magnitude of w .

In all of these plots, $P\{y = 1|x\} = 0.5$ when $x = 0$. This is a consequence of the fact that we used $b = 0$ for these plots.

Below is an interactive demonstration of these concepts, where you can select any values of b and w . By adjusting b , you can see the curve shift from left to right, and by adjusting w , you can see the transition sharpen.

By writing $z = b + wx = w(\frac{b}{w} + x)$ and recognizing that $\Pr\{y = 1|x\} = 0.5$ whenever $z = 0$, we can see that $\Pr\{y = 1|x\} = 0.5$ whenever $x = -\frac{b}{w}$.

```
In [14]: from ipywidgets import interact

def update_w(b, w):
    xmax = 5
    x = np.linspace(-xmax, xmax, 100)
    py = 1/(1+np.exp(-(b+w*x)))
    plt.plot(x, py)
    plt.grid()
    plt.title('b={0}, w={1}'.format(b, w))
    plt.xlabel('x')
    plt.ylabel('Pr{y=1|x}')
    plt.axis([-xmax, xmax, -0.1, 1.1])

interact(update_w, w = (-9, 11), b = (-10, 10));

interactive(children=(IntSlider(value=0, description='b', max=10, min=-10),
IntSlider(value=1, description='w' ...
```

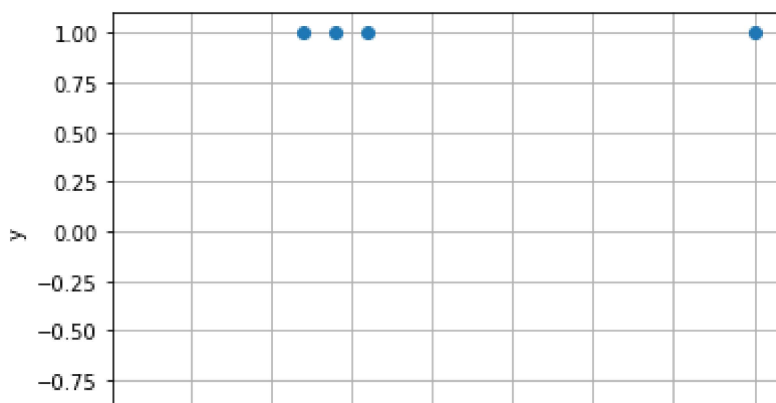
Demonstration on a Very Simple Dataset

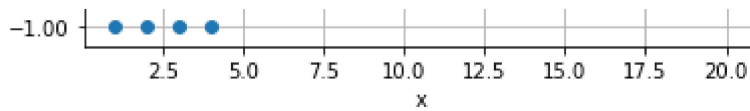
Next we use a simple data-fitting problem to illustrate the advantages of logistic regression over least-squares (LS) linear regression for binary classification.

Suppose we are given the eight (x_i, y_i) samples shown in the scatter plot below.

```
In [5]: xvec = np.array([ 1, 2, 3, 4, 6, 7, 8, 20])
yvec = np.array([-1, -1, -1, -1, 1, 1, 1, 1])

plt.plot(xvec, yvec, 'o')
plt.grid()
plt.xlabel('x')
plt.ylabel('y');
```





Just by looking at the dataset, we can see that a good rule would be to predict $y = 1$ whenever $x > 5$ and $y = -1$ when $x < 5$.

Try LS Linear Regression

Let's first try fitting a LS linear regression model to the data, to see if it behaves as expected.

Now, LS linear regression will return a real-number prediction, not a binary prediction as we desire for classification. To address this issue, we quantize the LS linear prediction (call it \hat{z}) to $\hat{y} = 1$ whenever $\hat{z} \geq 0$ and to $\hat{y} = -1$ whenever $\hat{z} < 0$.

Below, we fit the LS-linear model and plot \hat{z} and \hat{y} as a function of a test feature x .

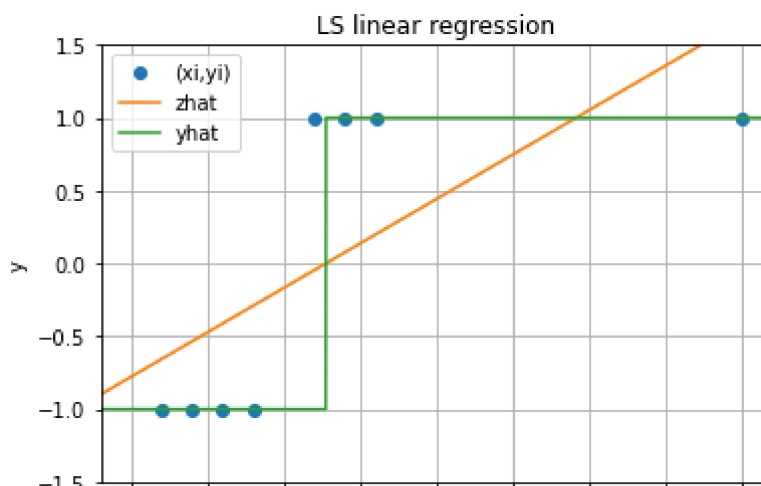
```
In [6]: from sklearn import datasets, linear_model
```

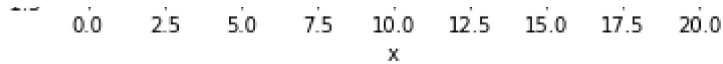
```
linreg = linear_model.LinearRegression()
linreg.fit(xvec.reshape(-1, 1), yvec)
b = float(linreg.intercept_)
w = float(linreg.coef_)
print('b=', b)
print('w=', w)
```

```
b= -0.7784342688330871
w= 0.12210733628754307
```

```
In [15]: xmin = 0
xmax = 20
x = np.arange(xmin-1, xmax+1, 0.01)
zhat = linreg.predict(x.reshape(-1, 1))
yhat = 2*np.array(zhat >= 0) - 1

plt.plot(xvec, yvec, 'o', label='(xi, yi)')
plt.plot(x, zhat, label='zhat')
plt.plot(x, yhat, label='yhat')
plt.grid()
plt.xlabel('x')
plt.ylabel('y')
plt.title('LS linear regression')
plt.legend()
plt.axis([xmin-1, xmax+1, -1.5, 1.5]);
```





The plot above shows that LS linear regression works to some extent, but not perfectly. It predicts $y = 1$ whenever $x > 6.25$ and $y = -1$ when $x < 6.25$, and thus mis-labels the $(x_i, y_i) = (6, 1)$ training sample.

What is happening is that the LS regression line (i.e., $\hat{z}(x) = \hat{b} + \hat{w}x$) is pulled to the right in order to fit the $(x_i, y_i) = (20, 1)$ sample. As a result, the zero-crossing of $\hat{z}(x)$ occurs at $x = 6.25$ and not at $x = 5$ as desired. This behavior will result whenever there is an imbalance in feature values of the two classes.

Try Logistic Regression

Next we will fit a logistic regression model to the dataset. The `sklearn` module provides a `LogisticRegression` method to do this. The interface is almost identical to the `LinearRegression` method that we used above.

Note that `sklearn`'s `LogisticRegression` uses L2 regularization by default. We don't need regularization for this example, so we'll disable it by setting `penalty=None`.

```
In [16]: logreg = linear_model.LogisticRegression(penalty='none')
```

```
logreg = linear_model.LogisticRegression()
logreg.fit(xvec.reshape(-1, 1), yvec)
b = float(logreg.intercept_)
w = float(logreg.coef_)
print('b=', b)
print('w=', w)
```

```
b= -5.62241853945352
w= 1.1208791859700278
```

```
In [17]: yhat = logreg.predict(x.reshape(-1, 1))
zhat = b + w*x
```

```
plt.plot(xvec, yvec, 'o', label='(xi, yi)')
plt.plot(x, zhat, label='zhat')
plt.plot(x, yhat, label='yhat')
plt.plot(x, f(zhat), label='Pr{y=1|x}')
plt.grid()
```